

Constraint driven Scheduling of fine-grained C concurrency for Reconfigurable Hardware

Akshay Gopalakrishnan
McGill University
akshay.akshay@mail.mcgill.ca

1 Introduction

The recent decades have seen the growing utilization of specialized hardware for the purpose of extensive computing. Right from general purpose CPUs to highly specialized hardware such as ASICs are being used in a wide variety of applications. When it comes to specialized hardware there is an increasing usage of FPGAs especially due to its advantage of reprogrammability as per user's need. In this endeavor, to synthesize such FPGA, high level synthesis(HLS) tools have been introduced. HLS is a process of High Level Synthesis(HLS) is a process of mapping behavioral(C-like programs) models of hardware to low level(HDL-like programs) ones; which are then synthesized as specialized hardware(eg: FPGA, ASIC). HLS tools are used widespread in recent times due to the high demand for such specialized hardware than have better throughput and energy consumption for specific intensive computations(eg: in Machine Learning).

However, HLS has been predominantly used/developed for sequential programs. The last few years have been investigated in terms of synthesizing concurrent programs; how would specialized hardware be synthesized that reflects concurrency. However, not much investigation has been done when it comes to synthesis of fine-grained concurrent programs; commonly known as lock-free programs.

2 Related Work

Several HLS tools are available today; some of the notables being Quartus and LegUp. LegUp allows their users to insert behavioral models of hardware in C language. Several works were done further to extend LegUp in synthesizing concurrent C programs to hardware.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

COMP 764, Winter 2022, Montreal, Quebec, Canada

© 2022 Association for Computing Machinery.

[1] propose a method to map software threads to parallel hardware for FPGAs. They do this for programs that utilize pthreads and OpenMP constructs. However, their focus was mainly on lock-based concurrency in C. They do note in their experiments that having a constraint on resources to map threads does effect the schedule(more cycles). Their contributions were integrated into LegUp.

[8] investigate the synthesis of concurrent programs utilizing fine grained concurrency elements of C. They show that on using atomics primitives provided by C, incorrect schedules can be obtained. They remedy this by adding additional ordering constraints among memory accesses. They test their approach by extending LegUp to tackle fine grained concurrent constructs of C. [9] is a follow up to this work implementing pipelined scheduling for such programs.

When it comes to HLS tools, one major advantage is the optimizations we can do at the behavioral model itself. These optimizations, have also been proven to significantly improve the resultant hardware synthesized. [2] and [3] show the impact/effect of compiler optimizations on HLS for Reconfigurable hardware like FPGAs. Specifically, they show how scheduling(termed as latency) has a major impact due the optimization passes of the compiler, as well as the order in which we perform it. But their work does not address this impact on concurrent programs(let alone those using fine grained Concurrency) synthesized to hardware.

While we do want to perform aggressive optimizations for concurrent behavioral models, we are constrained by the semantics of such programs. [10] shows that programs utilizing fine grained concurrency of C do not get the benefit of some optimizations as they are rendered unsafe in a concurrent context. [6] try to identify which memory accesses can be safely reordered in a weakly consistent C program targeted for hardware synthesis. Their goal was to improve the overall schedule of the program, and in turn the hardware that is synthesized for it. [7] is a follow up to this work describing a global analysis to more aggressively perform safe reorderings of memory accesses for the same purpose.

Previous works in the line of synthesizing weakly consistent C programs assume that every thread can be mapped to a unique hardware accelerator during synthesis. They do not investigate the effect of constraining the

number of available accelerators. While reordering was proposed in this line to improve schedules, no investigation was done on whether any optimization could have an impact on number of resources required for synthesis. Moreover, minimization of resources was not a factor considered in this effort; as the primary goal was to get overall correct schedule; and doing optimizations to help lower the required clock cycles.

3 Methodology

Previous work on synthesizing fine grained concurrent programs assumed infinite resources (hardware accelerators). In practice this is not true. One straightforward direction is to consider a more global scheduling approach to reuse resources. However, coming up with a sound global analysis that respects all consistency rules of the concurrency model might be a non-trivial task. Instead, I propose to address the resource constraint problem by performing an optimization that merges two or more threads. Previous work has shown that optimizations do significantly improve the resource requirements. However, the optimizations considered were only for sequential programs; while merging two threads is something for concurrent programs. Hence, we investigate whether this optimization indeed helps our case.

The main goal is to identify if we can minimize resource utilization while performing synthesis of concurrent programs such as above. We assume from previous work that synthesis is done per-thread; each thread is scheduled independently and allocated its own share of resources. In this endeavor, I have only addressed the correct scheduling and performing this optimization at the pre-scheduling stage. I do not implement the full HLS process; the program, if concurrent, cannot be synthesized into VHDL as of now.

3.1 Concurrency Model

There are several shared memory concurrency models that exist to date. Previous work has mainly focused on C based concurrency. The C concurrency model semantics is fairly non-trivial to tackle as well as implement. Hence, I choose a standard intuitive model widely understood; Sequential Consistency[4].

3.2 Types of Programs

The programs are considered to be a set of statements of the following format

$$a = b + cs \quad (1)$$

Here a refers to some memory. b and cs can be either memory or constants.

Memory can be termed as local (a and b) or global/shared (cs - any memory name ending with s).

```
gnode** ord_mem_list(basic_block *);
gnode** ord_mem_list_stmt(gnode **, gnode *, int *);
void set_sc_order(gnode **);
```

Figure 1. Three functions created to take a thread and add additional memory dependencies respecting Sequential Consistency.

```
//Check if all previous memory dependencies have already been scheduled
for(int i = 0; list[0]->mem_dep[i] != NULL; i++)
if(list[0] != NULL && list[0]->sched == -1) {
    if(list[0] -> mem_dep != NULL) {
        for(int i = 0; list[0]->mem_dep[i] != NULL; i++) {
            //printf("%d \n", i);
            if(list[0]->mem_dep[i]->sched == -1)
                return 0;
        }
    }
    list[0]->sched = clk_cycle;
    return 1;
}
```

Figure 2. Code snippet added to factor in memory dependencies while scheduling

Each program is viewed as a set of threads; each of which have the following statements as above. There are no conditional branching statements, nor loops.

3.3 Code Base Utilized

For the implementation, I reuse the existing code base from Assignment. The assignment dealt only with sequential programs of the above form. The extension thus, must be done to incorporate concurrency reflecting SC.

3.4 Additions to Code base

Threads. The notion of threads is easily captured by the code-base by representing each thread as its own DFG.

Correct Scheduling. One important thing to address is that of ensuring correct scheduling in a concurrent context. [8] showed that without adding extra memory dependencies, one can get an incorrect schedule for synthesis. Thus, the solution was to add additional dependencies between memory accesses which are shared. For sequential consistency, the dependencies were identified to be between every shared memory access in the same thread. Fig 1 represents the following functions that were created for this

Adding these dependency edges, the scheduler must also incorporate this information while allotting an action to a specific clock cycle time stamp. Fig 2 showcases the code snippet that was added for this. Note that because we still use the previous approach of independently scheduling each thread, these changes were identified to be sufficient for correct scheduling.

Merging Optimization. I have implemented a DFG as a set of DFGs in my code base. Thus, merging was

```
//Merge 2 t1-t2-t1-t2
merge[1]->dfts[0] = t1->dfts[0];
merge[1]->dfts[1] = t2->dfts[0];
merge[1]->dfts[2] = t1->dfts[1];
merge[1]->dfts[3] = t2->dfts[1];
```

Figure 3. Merging is equivalent to just adding and mixing DFG components of one thread with another

equivalent to simply adding one DFG into another. While merging, memory dependencies that are newly introduced must also be taken care of. This differentiates the different merging combinations. Since the DFG set is implemented as an array of DFGs, the implicit array order was used to retain the memory order after merging.

Fig 3 represents the above idea. Here $t1$ and $t2$ are two threads which are a set of DFGs. *merge* is the new program that is to be constructed using the elements of $t1$ and $t2$.

Soundness? An important factor to consider is whether such merging is a sound Transformation. [5], in their survey paper, do however, mention merging as a sound transformation under Sequential Consistency (SC). So a formal proof is not required.

4 Experimental Setup

I specifically wanted to capture different program patterns that captures the following features

- Captures Read-Write, Read-Read, Write-Read, Write-Write dependencies between shared memory.
- Logical resources (adders) used by such programs.
- Mixture of shared and local memory usage.

Since I assume my programs have no conditionals or loops, I wanted to capture relatively general program patterns that may possibly occur while performing optimizations. To capture all the above features, the following program format deemed sufficient as that in Fig 4.

$a1 = c1 + d1;$
 $b1 = e1 + f1;$

$a2 = c2 + d2;$
 $b2 = e2 + f2;$

Figure 4. Program Format

For each of the above accesses, I create a program where either they are shared or not. For example, $a1$ can be either just $a1$ or as (in this way, I capture the notion of shared memory between the two threads, since $a2$ may also become as .) This way, in total, I get 4096 programs.

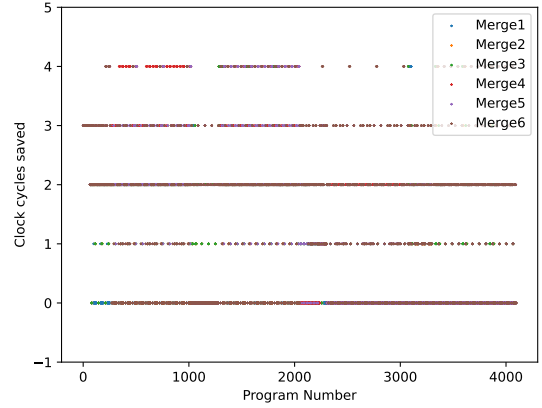


Figure 5. Scatter plot for each merging strategy, showing how many clock cycles were saved

5 Evaluation

The above programs were run and tested for each merging strategy. On merging, they were measured for the following:

- Clock cycles utilized after merging.
- Max resources that we require after merging.

I also measured before merging, for each program, the clock cycles and adders required for each thread. It is important to note that the total clock cycles represent the worst case scenario; if this program was run without merging.

The overall result was positive; it was possible to reduce or maintain the number of resources required, in addition to reducing the worst case number of clock cycles. Fig 5 showcases the clock cycles saved from the worst case on merging:

I noticed that for each program, the clock cycles saved is at least zero. This shows us that at least merging does not result in total clock cycles required to be more than the original program. The reason for this is quite intuitive; merging is equivalent to saying one thread executes first, followed by the second. Hence, naturally, the clock cycles will be bounded by the worst case before merging.

What is interesting however, is that the clock cycles are better than the worst case in almost 70percent of the programs that were considered after merging. Fig 6 shows the count of programs where (irrespective of how we merge) we save clock cycles on merge.

This shows us a promise that merging indeed may not be detrimental to number of clock cycles; there will exist some merging that helps us save clock cycles in most cases.

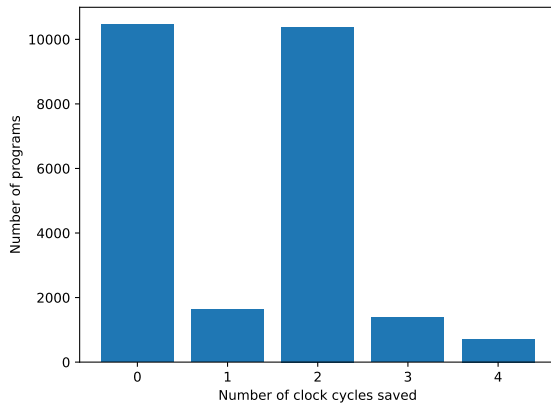


Figure 6. Count of programs after merging for the number of clock cycles saved.




Figure 7. Scatter plot for each merging strategy, showing how many adders were saved

When it comes to resources, Fig 7 showcases the resources saved (in comparison to pre-merge the amount of resources required). I noticed with the scatter plot that, once again, the amount of resources required is bounded by the amount required pre-merge. What is interesting however, is that resources are saved, even though the scheduling algorithm is greedy, assuming infinite resources. This occurred as adding additional memory dependency edges implicitly added dependency edges between the addition action too (this occurs due to Read-Write or Write-Read dependency edges).

Fig 8 shows us the count of programs (irrespective of how we merge) for the amount of resources we saved on merging. Once again, for most cases, some form of merging saved resources. This is promising; as this shows

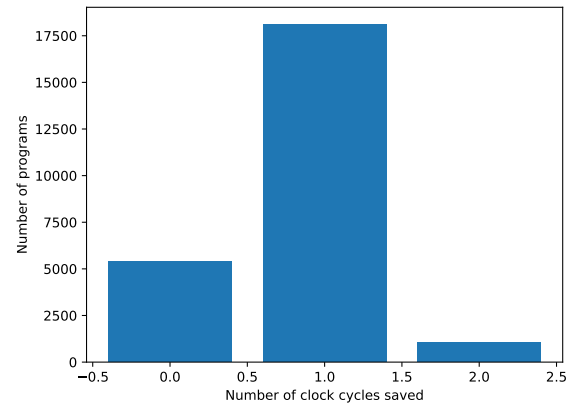


Figure 8. Count of programs after merging for the number adders saved.

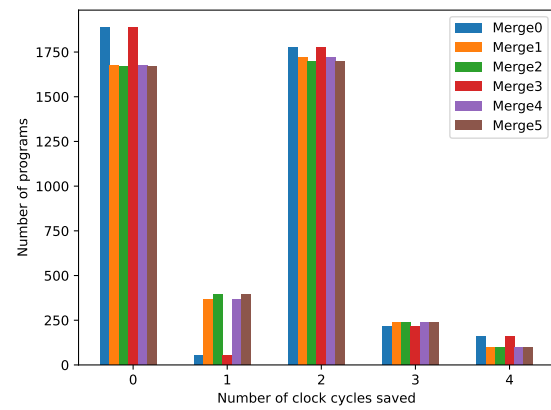


Figure 9. Graph showing for each merging strategy, the number of programs where merging saved clock cycles.

us that resource constraints can be addressed by an optimization which just merges threads.

Our testing was done with 6 possible mergings; which interleave program statements. Normally, merging is done by sequentializing a thread after another. But here, merging was performed on a more fine grained level. Is this useful though? Can we do away with just simple thread level merging? Turns out, interleaving forms of merging do help us; more clock cycles as well as more resources can be saved.

Fig 9 showcases for each merging strategy, the amount of programs that benefited in saving clock cycles. We note that the trend is same, overall around 70percent of the programs benefit by merging; saving clock cycles.

Fig 10 showcases for each merging strategy, the amount of programs that benefited in saving adders. We once




Figure 10. Graph showing for each merging strategy, the number of programs where merging saved adders.



Figure 11. Program for which Merging by interleaving program statements results in a better schedule.

again, note that in most cases, the amount of adders required after merging reduces at least by 1; irrespective of what merging strategy we use.

The above results show us that merging as an optimization is helpful in addressing the resource constraint problem. In addition, we do not detriment the overall clock cycles required for the program; this is always bounded by the worst case clock cycles of the original program. Moreover, in most cases, we even improve on this bound.

So far the above statistics does not indicate whether we require all forms of merging. It is important to note that simple merging one thread after the other did not always result in the best resource usage and clock cycles. For instance, for the program in Fig 11, the best merging optimization was interleaving t2's first statement after t1's first statement. Followed by t1's second statement and t2's second statement.

The schedule for merging this way is given in Fig 12. The blue arrow represents the memory dependency edges. The red circles represent the appropriate clock cycles when an action can be performed; as prescribed by greedy scheduling. The green boxes represent memory reads and blue boxes memory writes. The yellow arrow shows the order in which the two threads were merged. The overall schedule saved 2 clock cycles and 1 adder.

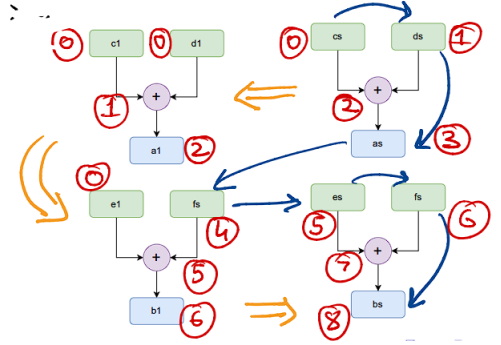


Figure 12. The schedule for merging in an interleaved fashion, giving us better schedule in comparison.

6 Discussion

The primary limitation of merging is that this optimization is still not sound for programming languages such as Java and C11. While behavioral models do intend to be written in C-style (for example in LegUp), one would need a different concurrency model to incorporate this optimization. The second limitation is the lack of conditionals and loops. Incorporating these would make merging at the implementation level more complex, as we would have to consider control flow also and merge parts of another thread only outside conditionals and loops of the primary thread, and so on. Lastly, the scheduling strategy tested for this has only been Greedy.

Future work in this line can be done by making merging more fine grained, and interleave at the level of memory accesses, and not program statements. The ideal goal would be to come up with a global analysis that identifies for a given concurrent program, the best merging optimization for sets of threads that save resources as well as clock cycles. In addition, it would be nice to see how other optimizations done by HLS tools play in conjunction with merging to give a better synthesized tool.

7 Conclusions

I tried to address the scheduling problem of fine grained concurrent programs. Previous work have placed no resource constraint while addressing this problem. I viewed the resource constraint problem as an optimization that can be done at the behavioral model level before scheduling. The optimization helped in addressing the resource constraint problem on most programs considered for benchmarking. It was noted that different variants of the optimization may be required to get least required resources.

References

- [1] Jongsok Choi, Stephen Dean Brown, and Jason Helge Anderson. 2013. From software threads to parallel hardware in high-level synthesis for FPGAs. In *2013 International Conference on Field-Programmable Technology, FPT 2013, Kyoto, Japan, December 9-11, 2013*. IEEE, 270–277. <https://doi.org/10.1109/FPT.2013.6718365>
- [2] Jason Cong, Bin Liu, Raghu Prabhakar, and Peng Zhang. 2012. A Study on the Impact of Compiler Optimizations on High-Level Synthesis. In *Languages and Compilers for Parallel Computing, 25th International Workshop, LCPC 2012, Tokyo, Japan, September 11-13, 2012, Revised Selected Papers (Lecture Notes in Computer Science)*, Hironori Kasahara and Keiji Kimura (Eds.), Vol. 7760. Springer, 143–157. https://doi.org/10.1007/978-3-642-37658-0_10
- [3] Qijing Huang, Ruolong Lian, Andrew Canis, Jongsok Choi, Ryan Xi, Stephen Dean Brown, and Jason Helge Anderson. 2013. The Effect of Compiler Optimizations on High-Level Synthesis for FPGAs. In *21st IEEE Annual International Symposium on Field-Programmable Custom Computing Machines, FCCM 2013, Seattle, WA, USA, April 28-30, 2013*. IEEE Computer Society, 89–96. <https://doi.org/10.1109/FCCM.2013.50>
- [4] Leslie Lamport. 1979. How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs. *IEEE Trans. Computers* 28, 9 (1979), 690–691. <https://doi.org/10.1109/TC.1979.1675439>
- [5] Evgenii Moiseenko, Anton Podkopaev, and Dmitrij V. Koznov. 2021. A Survey of Programming Language Memory Models. *Program. Comput. Softw.* 47, 6 (2021), 439–456. <https://doi.org/10.1134/S0361768821060050>
- [6] Nadesh Ramanathan, George A. Constantinides, and John Wickerson. 2018. Concurrency-Aware Thread Scheduling for High-Level Synthesis. In *26th IEEE Annual International Symposium on Field-Programmable Custom Computing Machines, FCCM 2018, Boulder, CO, USA, April 29 - May 1, 2018*. IEEE Computer Society, 101–108. <https://doi.org/10.1109/FCCM.2018.00025>
- [7] Nadesh Ramanathan, George A. Constantinides, and John Wickerson. 2021. Global Analysis of C Concurrency in High-Level Synthesis. *IEEE Trans. Very Large Scale Integr. Syst.* 29, 1 (2021), 24–37. <https://doi.org/10.1109/TVLSI.2020.3026112>
- [8] Nadesh Ramanathan, Shane T. Fleming, John Wickerson, and George A. Constantinides. 2017. Hardware Synthesis of Weakly Consistent C Concurrency. In *Proceedings of the 2017 ACM/SIGDA International Symposium on Field-Programmable Gate Arrays, FPGA 2017, Monterey, CA, USA, February 22-24, 2017*, Jonathan W. Greene and Jason Helge Anderson (Eds.). ACM, 169–178. <http://dl.acm.org/citation.cfm?id=3021733>
- [9] Nadesh Ramanathan, John Wickerson, and George A. Constantinides. 2018. Scheduling Weakly Consistent C Concurrency for Reconfigurable Hardware. *IEEE Trans. Computers* 67, 7 (2018), 992–1006. <https://doi.org/10.1109/TC.2017.2786249>
- [10] Viktor Vafeiadis, Thibaut Balabonski, Soham Chakraborty, Robin Morisset, and Francesco Zappa Nardelli. 2015. Common Compiler Optimisations are Invalid in the C11 Memory Model and what we can do about it. In *Proceedings of the 42nd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL 2015, Mumbai, India, January 15-17, 2015*, Sriram K. Rajamani and David Walker (Eds.). ACM, 209–220. <https://doi.org/10.1145/2676726.2676995>